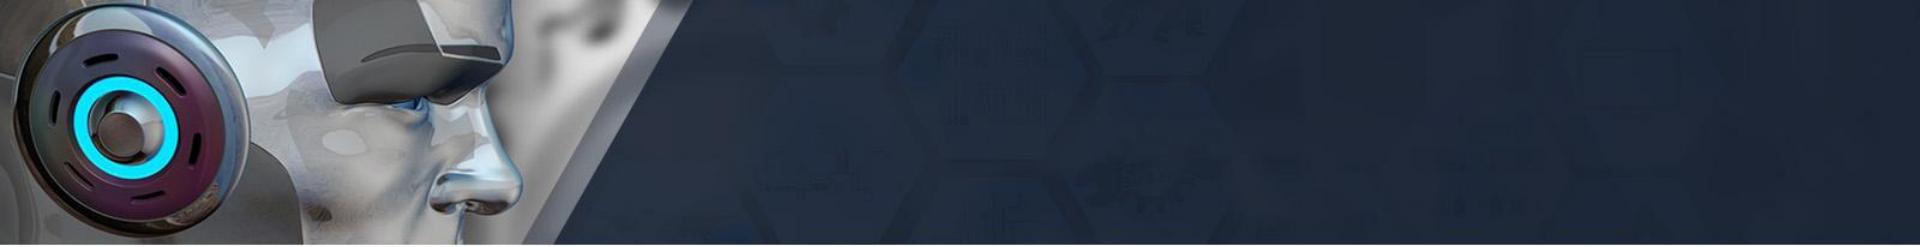




LLM Compression: How Far Can We Go in Balancing Size and Performance?

Dr. Shantipriya Parida
Senior AI Scientist

AMD Silo AI, Helsinki, Finland



Agenda

- Motivation & Challenges
- Compression Methods
- Use Cases
- Takeaways



Motivation and Challenges

- LLMs provide powerful capabilities but require high memory and compute
- Resource constraints hinder deployment on edge and mobile devices
- Model compression is essential for real-world applicability
- Naïve low-bit quantization degrades model accuracy
- Need for accuracy-aware, efficient compression methods

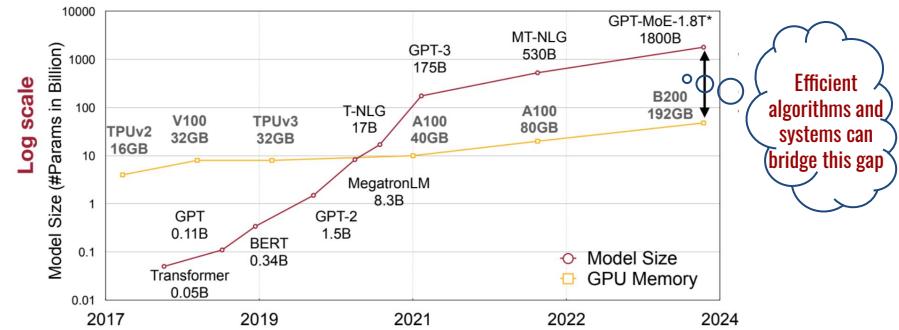


Fig 1: Exponential Growth of LLM Parameters Outpacing GPU Memory Capacity
Image source: [AWQ: Activation-aware Weight Quantization](#)

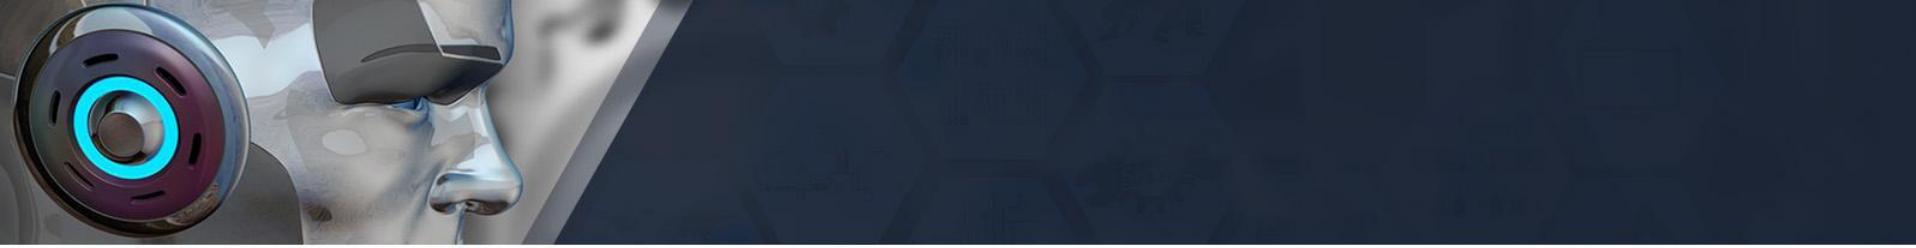




Motivation and Challenges

- **Edge Computing Applications**
 - Deploy sophisticated computer vision and NLP models directly on IoT devices and sensors, enabling real-time processing without cloud connectivity requirements.
- **Mobile AI Development**
 - Integrate complex deep learning capabilities into smartphone applications while maintaining battery efficiency and performance responsiveness.
- **Resource-Limited Research Environments**
 - Enable AI research teams with limited computational resources to work with state-of-the-art models through efficient compression techniques.





Model Compression Techniques

- **Quantization**

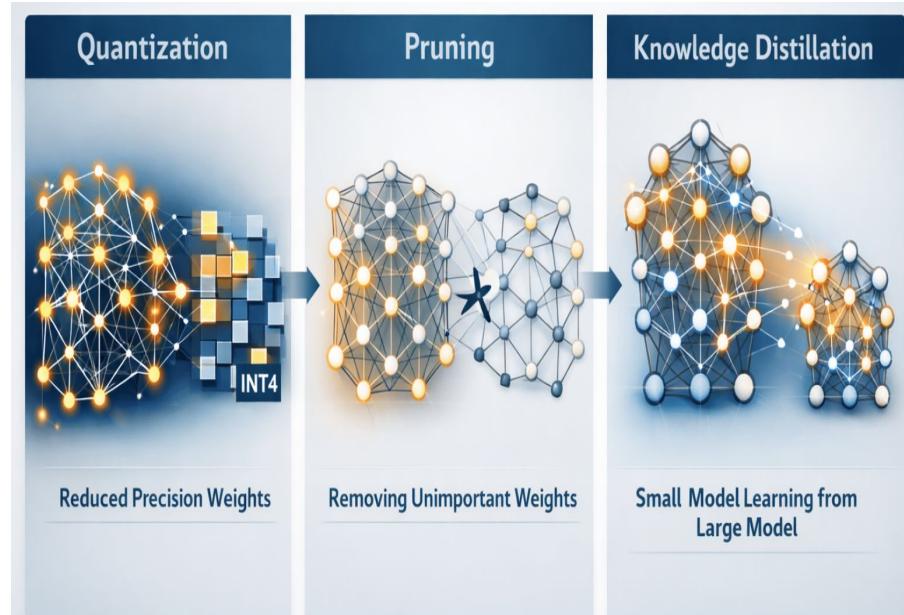
- Reduce numerical precision of model weights
- Lowers memory footprint and inference cost
- Enables deployment on resource-constrained hardware

- **Pruning**

- Remove redundant or less important weights
- Preserves performance while reducing model size
- Improves inference efficiency

- **Knowledge Distillation**

- Train a smaller model using a large teacher model
- Retains key behaviors of the original LLM
- Produces compact, high-performing models





What is Quantization ?

- Quantization aims to reduce the precision of a model's parameter from higher bit-widths (like 32-bit floating point) to lower bit-widths (like 8-bit integers).
- Reduce numerical precision of model weights from 32-bit floating point to lower bit representations (8-bit, 4-bit), dramatically decreasing memory requirements while maintaining functional accuracy.

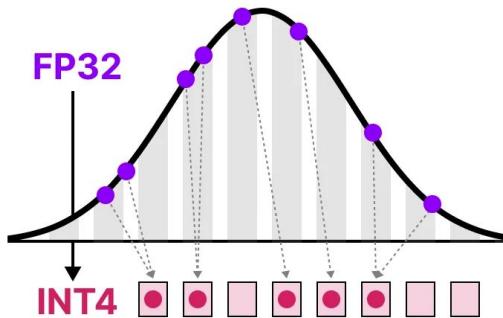


Image source:

<https://newsletter.maartengrootendorst.com/p/a-visual-guide-to-quantization>



Why Memory Matters in LLMs

- **Model size = memory footprint**
Number of parameters directly determines RAM/GPU memory needs (FP32/FP16 models quickly exceed hardware limits).
- **Precision defines memory efficiency**
Moving from FP32 → FP16 → INT8 → INT4 reduces memory by **2×–8×** with minimal accuracy loss.
- **Quantization is the key enabler**
Large models (e.g., 70B) become deployable on a **single GPU** when stored in 4-bit representations.
- **Memory vs performance trade-off**
Larger models offer better reasoning and accuracy, but smaller or quantized models deliver **higher throughput and lower latency**.
- **Perplexity ≠ real-world quality**
Slight increases in perplexity after compression often **do not hurt downstream task performance**.
- **Optimal choice is task-dependent**



Numerical Representations in LLMs

(IEEE-754)

Floating-Point Formats

FP32 (Single Precision – 32 bits)

- Sign: 1 bit, Exponent: 8 bits, Mantissa (Fraction): 23 bits
- High numerical precision and wide dynamic range
- High memory and compute cost

FP16 (Half Precision – 16 bits)

- Sign: 1 bit, Exponent: 5 bits, Mantissa (Fraction): 10 bits
- Lower precision and dynamic range
- Reduced memory footprint and faster computation

BF16 (BFloat16 – 16 bits)

- Sign: 1 bit | Exponent: 8 bits | Mantissa: 7 bits
- Maintains FP32-like dynamic range
- Slightly lower precision, ideal for stable training
- Widely adopted in TPU & GPU training

Integer & Low-Bit Formats (Used in LLM Compression)

INT8 (8 bits)

- Fixed-point integer (no exponent or mantissa)
- Typically uses per-tensor or per-channel scaling factors
- Significantly reduces memory (4x vs FP32)
- Widely used for post-training quantization and inference

INT4 (4 bits)

- Extremely low-bit integer representation
- Requires careful scaling, grouping (e.g., GPTQ, AWQ, GSQ)
- Provides major memory and bandwidth savings
- Popular for LLM deployment on GPUs and edge devices

INT2 / Binary (Research & Experimental)

- Ultra-low precision formats
- Aggressive compression with higher accuracy risk
- Mostly experimental for LLMs

Float 16-bit (FP16)



Fig 2: FP16 Example

Precision Formats Used in Modern LLM Pipelines

Format	Used Where	Why
32 FP32	🔍 Debugging, critical ops	🌐 Stability
16 FP16	🚀 Training & inference	⌚ Speed
BF16 BF16	GPU Training (TPUs, GPUs)	⚡ Wider exponent
8 INT8	📊 Inference	💲 Cost & latency
4 INT4	💻 Edge / large-scale serving	🏗️ Extreme compression

Fig 3: Precision Formats Used in Modern LLM Pipelines.



LLM Precision & Quantization

Value Representation Across Precisions

- Same value → different representations across precisions (Fig 4)
- **FP32**: high precision, high memory & compute
- **FP16**: faster, cheaper, lower numerical accuracy

Conversion to Lower Precision (Quantization)

- FP32 values mapped to discrete integer levels (Fig 5)
- Reduces model size and memory bandwidth
- Commonly applied during **inference**, not full training

Why This Matters ?

Enables larger batch sizes, reduces GPU memory footprint, lowers latency and serving cost, and makes LLMs deployable at scale

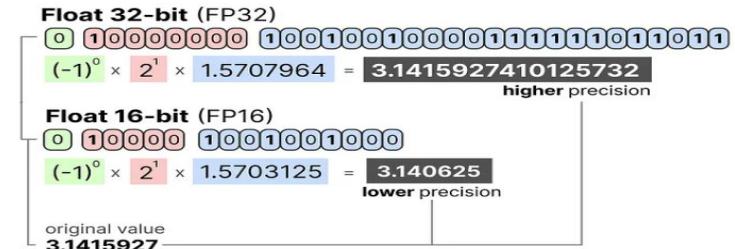


Fig 4: Floating-Point Precision Comparison (FP32 vs FP16). Reduced mantissa in FP16 leads to rounding error but faster computation

Image Source: <https://medium.com/data-science/a-visual-guide-to-quantization-930ebcd9be94>

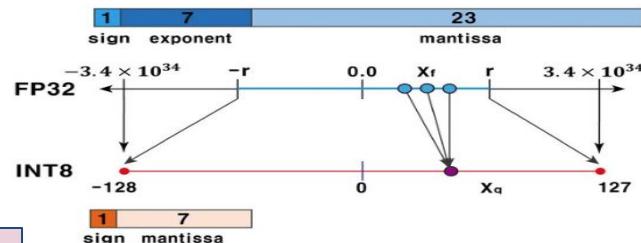


Fig 5: Quantization from FP32 to INT8.
Continuous FP32 values mapped to discrete integer levels

Image Source:
https://www.researchgate.net/publication/363946099_A_Method_of_Deep_Learning_Model_Optimization_for_Image_Classification_on_Edge_Device



Other Quantization Approaches – Post Training Quantization (PTQ)

Post-training quantization occurs when quantization is applied to an existing model. This converts the model from a floating-point representation to a lower precision fixed-point integer without the need for any retraining. This method does not require as much data as quantization aware training and is much faster.

Advantages:

- **Simplified training:** PTQ eliminates the need for modifications during model training, making it a more straightforward approach.
- **Flexibility:** PTQ allows for the use of pre-trained models, which can be beneficial when working with existing models or when the original training data is unavailable.

Challenges:

Accuracy loss: PTQ can introduce accuracy loss, particularly for smaller networks, due to the reduced precision of model weights and activations



Quantization Approaches - Quantization Aware Training (QAT)

- Quantization-aware training incorporates the conversion of weight during the pre-training or fine-tuning of an LLM.
- This allows for enhanced performance, but it demands a large amount of computational power and requires representative training data.
- QAT will usually produce a model with higher performance, but it is more expensive and will demand far more computing power.

Key Components

- **Fake Quantization:** Simulate quantization noise by applying quantization and dequantization operations to the model's weights and activations during training.
- **Quantization Configuration:** Specify the quantization parameters, such as bit-width, scaling, and zero-point, to control the quantization process.

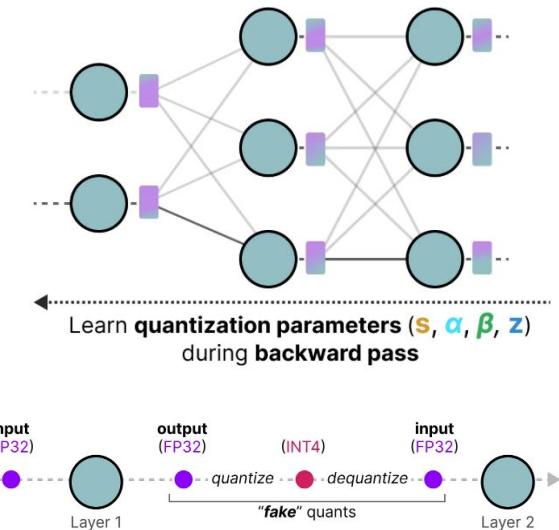


Fig 6: Quantization-Aware Training (QAT) Workflow



Quantization Steps

- **Weight Quantization:** Convert model weights from FP32 to a lower precision data type (e.g., INT8). This involves:
 - Determining the range of weight values
 - Mapping the weight values to a smaller set of discrete values (e.g., integers)
 - Minimizing information loss through techniques like: Linear quantization & Non-linear quantization (e.g., histogram-based)
- **Activation Quantization:** Convert model activations (output of each layer) from FP32 to a lower precision data type (e.g., INT8). This involves:
 - Estimating the range of activation values
 - Mapping the activation values to a smaller set of discrete values (e.g., integers)
 - Minimizing information loss through techniques like:
 - Linear quantization
 - Non-linear quantization (e.g., histogram-based)
 - Learned quantization (e.g., using neural networks)
- **Model Calibration:** Adjust the quantized model to minimize errors and maintain accuracy. This may involve:
 - Fine-tuning the model with the quantized weights and activations
 - Adjusting the quantization parameters (e.g., number of bits, quantization intervals)



Popular Quantization Techniques

Generative Pretrained Transformer Quantization (GPTQ)

- One-shot post-training quantization method for large language models
- Converts model weights from FP32/FP64 to low-bit precision (e.g., 8-bit or 4-bit)
- Performs row-wise weight quantization, optimizing each row independently
- Minimizes accuracy loss using second-order error minimization and dynamic error compensation
- Adjusts later weights to offset quantization errors introduced earlier
- Enables faster and more memory-efficient inference due to low-precision arithmetic

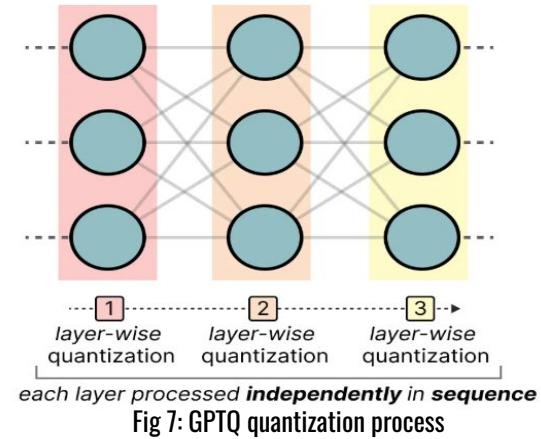


Fig 7: GPTQ quantization process





Popular Quantization Techniques

Group Scaling Quantization (GSQ)

- A post-training quantization method inspired by **Activation-Aware Quantization (AWQ)**
- **Divides weight matrices into groups**, assigning a **shared scaling factor** per group
- Ensures all quantized weights fit within the **INT4 range**, reducing precision loss
- **Selects important weights based on activation impact**, not just weight magnitude
- Preserves **fine-grained activation information**, improving post-quantization accuracy
- Particularly effective for **4-bit quantization** of LLMs

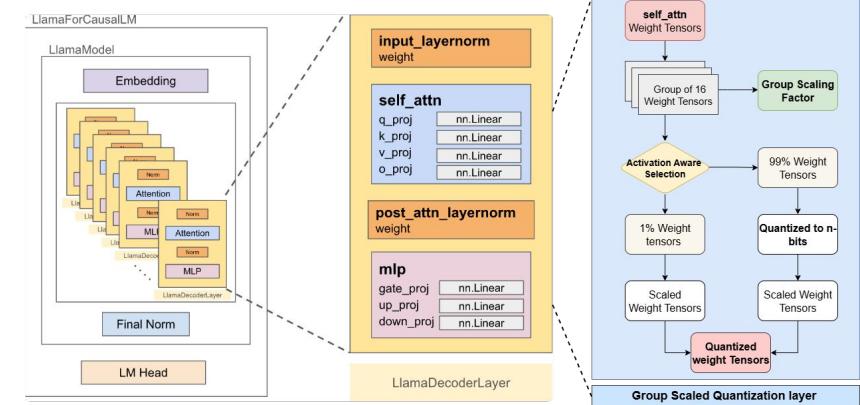
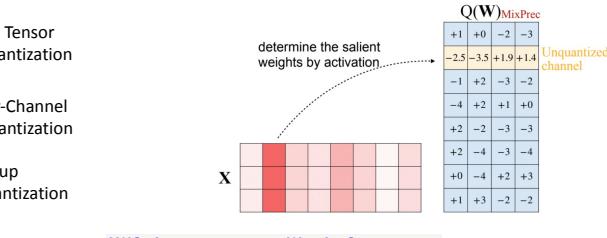
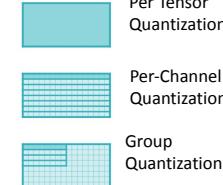


Fig: 8: Detailed view of the architecture behind GSQ quantization process. 14



Use Case

Objective: Evaluate GPTQ for post-training compression of small LLMs under realistic inference workloads.

Models Quantized

- LLaMA 1B
- Qwen 0.5B
- Phi 1.5B

Quantization Method

- GPTQ (Post-Training Quantization)
- Weight-only INT4 quantization
- One-shot calibration (no retraining)

Evaluation Tasks

- Information Retrieval
- Question Answering
- Mathematical Reasoning

Datasets

- MS MARCO (IR)
- BoolQ (QA)
- GSM8K (Math)

Metrics

- Accuracy, Perplexity, latency, throughput, Memory

Work Published: Sk, S., Dhal, D., Khosla, S., Dhaka, A., Parida, S., Shahid, S., ... & Bojar, O. (2025, September). LLM Compression: How Far Can We Go in Balancing Size and Performance?. In *Proceedings of the 15th International Conference on Recent Advances in Natural Language Processing-Natural Language Processing in the Generative AI Era* (pp. 1183-1187).



Use Case:

Key Results & Trade-offs

- **LLaMA-1B:** Quantization significantly improves MS MARCO ($81.1 \rightarrow 99.9\%$) and BoolQ ($40.1 \rightarrow 62.2\%$) with negligible impact on GSM8K
- **Qwen-0.5B:** More sensitive to quantization; MS MARCO drops, while BoolQ remains stable
- **Phi-1.5B:** Highly robust; performance remains **almost unchanged** after quantization
- **Retrieval & QA tasks tolerate low-bit quantization better than reasoning tasks (GSM8K)**
- **Larger models** show greater resilience to aggressive quantization

Trade-offs

- Up to **13x model size reduction** with minor accuracy loss
- Structured quantization preserves critical parameters
- Slight latency increase from group-level scaling

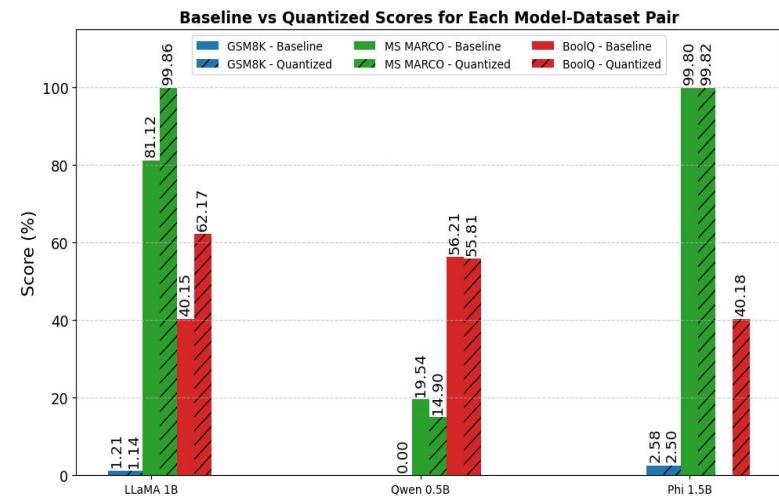
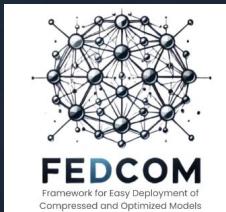


Fig: 9: Full-Precision vs. Quantized Model Performance



Takeaways

- Quantization enables deployment of LLMs on resource-limited hardware.
- **GPTQ** and **GSQ** provide different accuracy–efficiency trade-offs.
- Activation-aware and structured quantization preserve critical model behaviors.
- Model size can be reduced **up to 13 \times** with minimal performance loss.
- Quantization strategy should be **task-dependent** (e.g., reasoning vs. retrieval).
- Careful selection of methods is crucial for balancing **accuracy, latency, and memory**.



What Next...

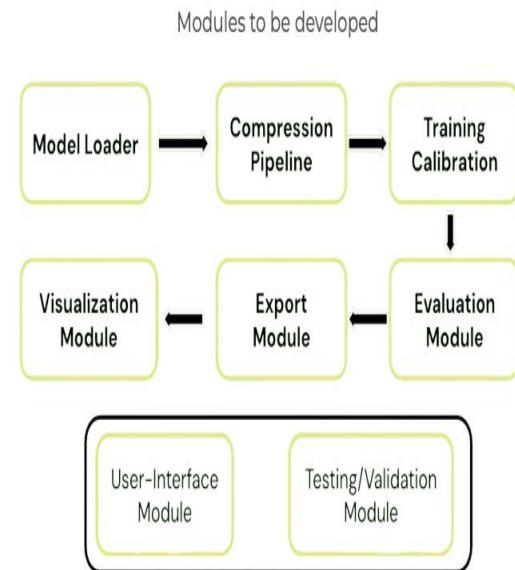
Framework for Easy Deployment of Compressed and Optimized Models (FEDCOM)

Goal: FEDCOM simplifies deep learning model compression and deployment by providing a structured workflow for quantization, pruning, distillation, and optional fine-tuning — making efficient models accessible for resource-constrained environments.

Features: It supports multimodal models, configurable compression techniques, performance evaluation (size, speed, accuracy), visual insights into trade-offs, and export of deployment-ready compressed models.

Purpose: Designed to reduce technical complexity and enable users of varied expertise to efficiently compress and deploy optimized models for real-world scenarios

High-Level Architecture



Open to research ideas and open-source contributions.



References

- Sk, S., et al. "LLM Compression: How Far Can We Go in Balancing Size and Performance?" *Proceedings of the International Conference on Recent Advances in Natural Language Processing (RANLP 2025)*, 2025, pp. 1183–1187. ACL Anthology, <https://aclanthology.org/2025.ranlp-1.136.pdf>.
- Frantar, Elias, et al. "GPTQ: Accurate Post-Training Quantization for Generative Pretrained Transformers." *arXiv*, 2024, <https://arxiv.org/abs/2210.17323>.
- Lin, Ji, et al. "Group Scaling Quantization (GSQ) for Large Language Models: Activation-Aware Quantization." *arXiv*, 2024, <https://arxiv.org/abs/2306.00978>.
- Lee, Byungsoo, et al. "LittleBit: Ultra Low-Bit Quantization via Latent Factorization." *arXiv*, 2025, <https://arxiv.org/abs/2506.13771>.
- Grootendorst, Maarten. "A Visual Guide to Quantization." *Maarten Grootendorst*, 2024, <https://maartengrootendorst.com>.
- Grootendorst, Maarten. "Quantization in vLLM: From Zero to Hero." *Maarten Grootendorst*, 2024, <https://maartengrootendorst.com>.
- Grootendorst, Maarten. "Quantization Notes." *Maarten Grootendorst*, 2024, <https://maartengrootendorst.com>.
- Grootendorst, Maarten. "Optimizing LLMs: Does Size Really Matter?" *Maarten Grootendorst*, 2024, <https://maartengrootendorst.com>.
- Grootendorst, Maarten. "How to Quantize an LLM with GGUF or AWQ." *Maarten Grootendorst*, 2024, <https://maartengrootendorst.com>.



Acknowledgement

Acknowledgement is given to **Sahil Sk, Debasish Dhal, Sonal Khosla, Sk Shahid, Sambit Shekhar, Akash Dhaka, Shantipriya Parida, Dilip K. Prasad, and Ondřej Bojar** for their work titled "**LLM Compression: How Far Can We Go in Balancing Size and Performance?**"



Thank you

Scan to access the Presentation slides

